

bayesian-bandits-code

March 15, 2020

```
[55]: import matplotlib.pyplot as plt
import numpy as np
from scipy.stats import beta, uniform
import seaborn as sns
```

0.1 Example 1

```
[36]: NUM_TRIALS = 2000
BANDIT_PROBABILITIES = [0.2, 0.5, 0.75]
```

```
[30]: class Bandit(object):
    def __init__(self, p):
        self.p = p
        self.a = 1
        self.b = 1

    def pull(self):
        return np.random.random() < self.p

    def sample(self):
        return np.random.beta(self.a, self.b)

    def update(self, x):
        self.a += x
        self.b += 1 - x

def plot(bandits, trial):
    x = np.linspace(0, 1, 200)
    for b in bandits:
        y = beta.pdf(x, b.a, b.b)
        plt.plot(x, y, label='real p: %.4f' % b.p)
    plt.title('Bandit distributions after %s trials' % trial)
    plt.legend()
    plt.show()
```

```

def experiment():
    bandits = [Bandit(p) for p in BANDIT_PROBABILITIES]

    sample_points = [5, 10, 20, 50, 100, 200, 500, 1000, 1500, 1999]

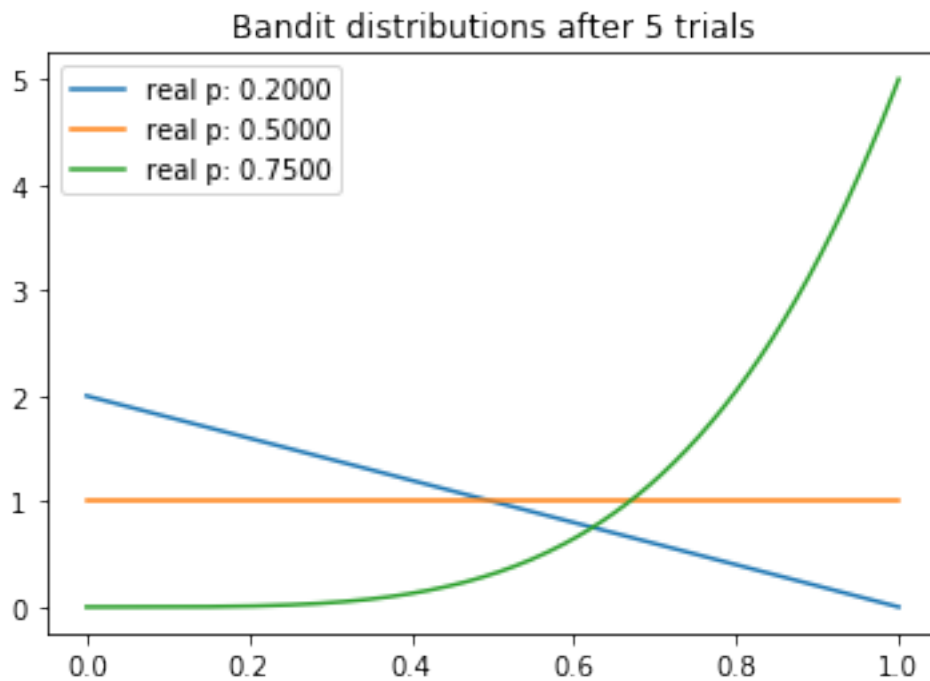
    for i in range(NUM_TRIALS):
        bestb = None
        maxsample = -1
        allsamples = []
        for b in bandits:
            sample = b.sample()
            allsamples.append('%.4f' % sample)
            if sample > maxsample:
                maxsample = sample
                bestb = b
        if i in sample_points:
            print('current samples: %s' % allsamples)
            plot(bandits, i)

    x = bestb.pull()
    bestb.update(x)


```

[31]: experiment();


current samples: ['0.7618', '0.4567', '0.8549']




current samples: ['0.1935', '0.0163', '0.7698']




current samples: ['0.0806', '0.5004', '0.9161']




current samples: ['0.2050', '0.1388', '0.8971']




current samples: ['0.0555', '0.0838', '0.8012']




current samples: ['0.4525', '0.0752', '0.7951']




current samples: ['0.0559', '0.0381', '0.7579']




current samples: ['0.1711', '0.3887', '0.7536']



current samples: ['0.1885', '0.3112', '0.7767']



current samples: ['0.1058', '0.1914', '0.7520']



0.2 Example 2

```
[59]: '''  
This function takes as input three tuples for alpha,beta that specify  
→priorR,priorG,priorB  
And returns R,G,B along with the maximum value sampled from these three  
→distributions.  
We can sample from a beta distribution using scipy.  
'''  
def find_asset(priorR,priorG,priorB):  
    red_rv = beta.rvs(priorR[0],priorR[1])  
    green_rv = beta.rvs(priorG[0],priorG[1])  
    blue_rv = beta.rvs(priorB[0],priorB[1])  
    return assets[np.argmax([red_rv,green_rv,blue_rv])]  
  
'''  
This is a helper function that simulates the real world using the actual  
→probability value of the assets.
```

*In real life we won't have this function and our user click input will be the
 → proxy for this function.*

```
'''
def simulate_real_website(asset, real_probs_dict):
    #simulate a coin toss with probability. Asset clicked or not.
    if real_probs_dict[asset]> uniform.rvs(0,1):
        return 1
    else:
        return 0
'''
```

*This function takes as input the selected asset and returns the posteriors for
 → the selected asset.*

```
'''
def update_posterior(asset,priorR,priorG,priorB,outcome):
    if asset=='R':
        priorR=(priorR[0]+outcome,priorR[1]+1-outcome)
    elif asset=='G':
        priorG=(priorG[0]+outcome,priorG[1]+1-outcome)
    elif asset=='B':
        priorB=(priorB[0]+outcome,priorB[1]+1-outcome)
    return priorR,priorG,priorB
'''
```

This function runs the strategy once.

```
'''
def run_strategy_once(priorR,priorG,priorB):
    # 1. get the asset
    asset = find_asset(priorR,priorG,priorB)
    # 2. get the outcome from the website/users
    outcome = simulate_real_website(asset, real_probs_dict)
    # 3. update prior based on outcome
    priorR,priorG,priorB = update_posterior(asset,priorR,priorG,priorB,outcome)
    return asset,priorR,priorG,priorB
'''
```

```
def plot_posteriors(priorR,priorG,priorB,ax=None,title=None):
    #fig = plt.figure(figsize=(12.5, 10))
    parameters = [priorR,priorG,priorB]
    x = np.linspace(0.001, 1, 150)
    for i, (_alpha, _beta) in enumerate(parameters):
        color = assets[i]
        y = beta.pdf(x, _alpha, _beta)
        lines = sns.lineplot(x, y, label="%s (%.1f,%.1f)" % (color, _alpha,
→ _beta), color = color,ax=ax)
        plt.fill_between(x, 0, y, alpha=0.2, color=color)
    if title:
```



```

        plt.title(title)
        plt.autoscale(tight=True)
        plt.legend(title=r"$\alpha$, $\beta$ - parameters")
        return plt

```

```

[60]: real_probs_dict = {'R':0.8, 'G':0.4, 'B':0.3}
      assets = ['R', 'G', 'B']

      priorR, priorG, priorB = (1,1), (1,1), (1,1)

      data = [('_', priorR, priorG, priorB)]

      for i in range(50):
          asset, priorR, priorG, priorB = run_strategy_once(priorR, priorG, priorB)
          data.append((asset, priorR, priorG, priorB))

```

```

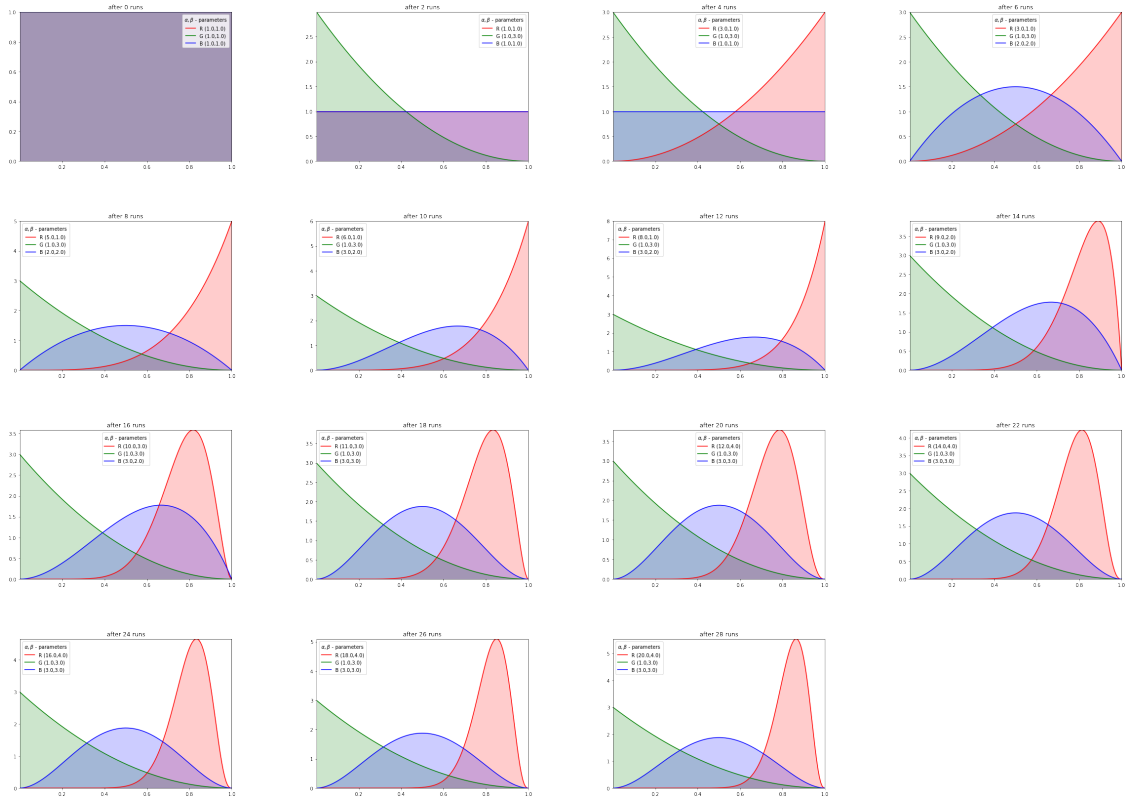
[61]: fig = plt.figure(figsize=(40, 60))
      fig.subplots_adjust(hspace=0.4, wspace=0.4)
      cnt=1
      for i in range(0, 30, 2):
          ax = fig.add_subplot(8, 4, cnt)
          g = plot_posteriors(*data[i][1:], ax, "after "+str(i)+" runs")
          cnt+=1
      plt.show()

```

```

/Users/jujohnson/anaconda3/envs/tf.latest/lib/python3.6/site-
packages/ipykernel_launcher.py:57: MatplotlibDeprecationWarning: Support for
uppercase single-letter colors is deprecated since Matplotlib 3.1 and will be
removed in 3.3; please use lowercase instead.
/Users/jujohnson/anaconda3/envs/tf.latest/lib/python3.6/site-
packages/IPython/core/pylabtools.py:128: MatplotlibDeprecationWarning: Support
for uppercase single-letter colors is deprecated since Matplotlib 3.1 and will
be removed in 3.3; please use lowercase instead.
    fig.canvas.print_figure(bytes_io, **kw)

```



[:

[: